

# Adversarial training for blind-spot removal (HiWi Report)

Rishi Sharma

February 11, 2020

# Outline

- 1 Adversarial Training
- 2 LSTM Network for DGA Prediction
- 3 Non-Differentiable Embedding Layer
- 4 Gradient-based Attack
- 5 Iterative Hardening
- 6 Transferring Adversaries
- 7 Future Work

## Natural saddle point (min-max) formulation

$$\min_{\theta} \rho(\theta), \text{ where } \rho(\theta) = \mathbb{E}_{(x,y) \sim D} [\max_{\delta \in S} L(\theta, x + \delta, y)]$$

- $\theta$  : the parameters of the model.
- $x$  : the input to the model.
- $y$  : the target label of the given input.
- $S \subseteq \mathbb{R}^n$  : the set of allowed perturbations.
- Two separate problems: Inner Maximization & Outer Minimization

[Towards Deep Learning Models Resistant to Adversarial Attacks (2017), Aleksander Madry et al.]

## Inner Maximization

$$\max_{\delta \in S} L(\theta, x + \delta, y)$$


- This optimization tries to generate the adversarial example from the given input.
- This optimization problem can be solved by using techniques such as Fast Gradient Sign Method and Projected Gradient Descent.
- FGSM:  $x \leftarrow x + \epsilon \text{sgn}(\nabla_x L(\theta, x, y))$ .
- PGD:  $x^{t+1} \leftarrow \Pi_{x+S}(x^t + \alpha \text{sgn}(\nabla_x L(\theta, x, y)))$ .

## Outer Minimization

$$\min_{\theta} \rho(\theta)$$

- This is the learning phase of the adversarial training.
- The network is trained on the dataset augmented with the adversarial examples.

# The Network

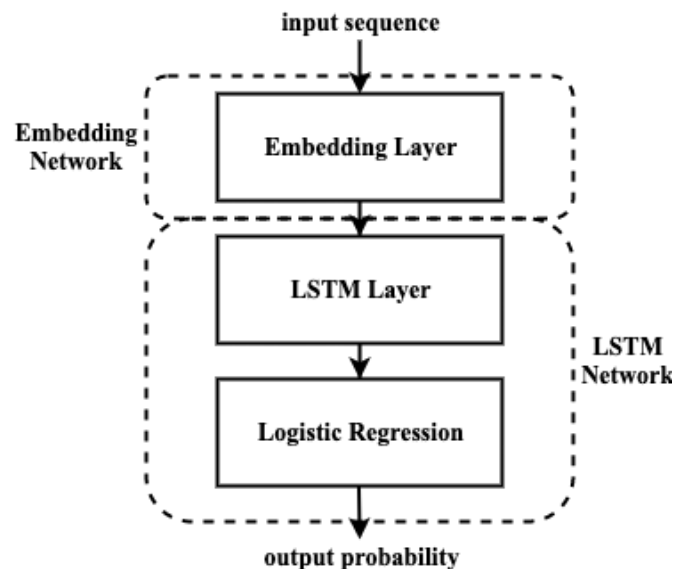


## Composition of the Network.

- The Embedding Layer : projects the  $l$ -length (padded) input sequence to a sequence of  $l$  vectors, each of dimension  $d$ .
- The LSTM layer : a feature extraction layer.
- Logistic Regression : predicts the probability of being malicious.

[Predicting Domain Generation Algorithms with Long Short-Term Memory Networks (2016), Endgame, Inc.]

# The Network



## Separation of the Networks.

- Trained over  $\sim 22k$  samples of each non-manipulated benign and malicious samples.
- 5-fold cross-validation.


-	Mean	Stddev
Acc	0.9974	0.0011
FNR	0.0014	0.0010
TNR	0.9961	0.0018
TPR	0.9986	0.0010
FPR	0.0039	0.0018

**Table 1:** The performance measurements of training the original network.

- This training was performed on only valid domain names. We will see this later why?




# Non-Differentiable Embedding Layer



A forward pass through the network.

# Non-Differentiable Embedding Layer



The gradient flow.

# Non-Differentiable Embedding Layer

## Problem

The Embedding Layer **selects** a vector corresponding each character in the input sequence.

- The implementations of the Embedding Layer in frameworks access the vector from a table using the character as the index.
- The layer is non-differentiable.

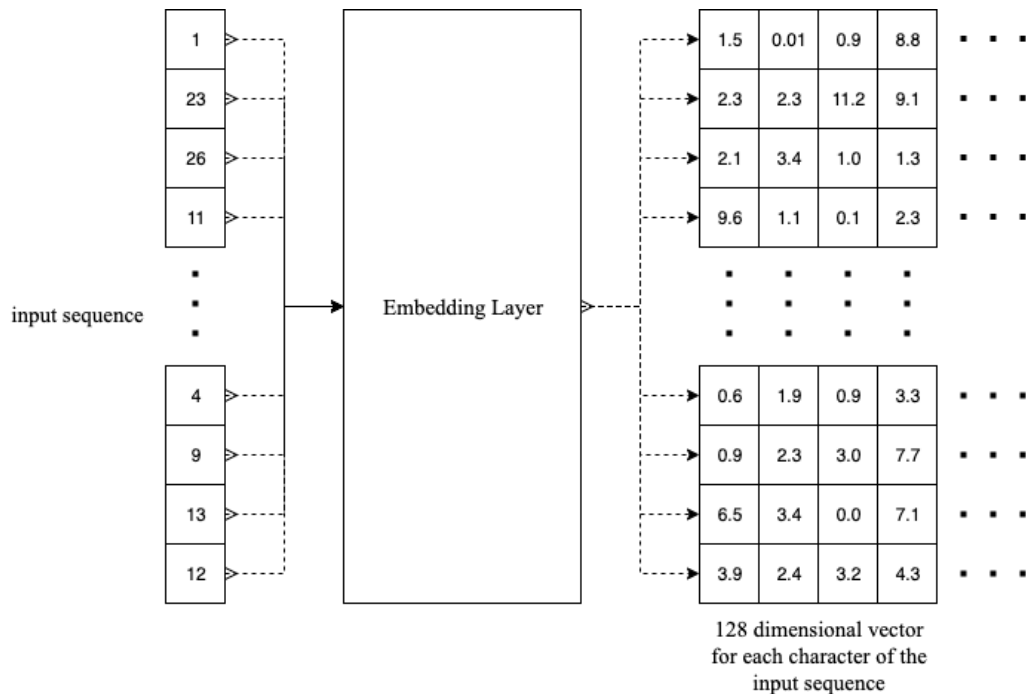
# Non-Differentiable Embedding Layer

## Problem

The Embedding Layer **selects** a vector corresponding each character in the input sequence.

- The implementations of the Embedding Layer in frameworks access the vector from a table using the character as the index.
- The layer is non-differentiable.
- Two solutions.

# Non-Differentiable Embedding Layer



The Embedding Network.

# Non-Differentiable Embedding Layer

## Emulate Embedding Layer

Use 1D Convolution filters to learn a vector representation of the input sequence elements.

- A very simple network.
- Each Convolution filter learn to predict one of the dimensions of the vector representation.
- This embedding layer is differentiable!

# Non-Differentiable Embedding Layer

## Emulate Embedding Layer

Use 1D Convolution filters to learn a vector representation of the input sequence elements.

- A very simple network.
- Each Convolution filter learn to predict one of the dimensions of the vector representation.
- This embedding layer is differentiable!
- *Con* : The accuracy fell down awfully on training the complete network in one pass.

# Non-Differentiable Embedding Layer

## Emulate Embedding Layer

Use 1D Convolution filters to learn a vector representation of the input sequence elements.

-	Mean	Stddev
Acc	0.9244	0.0388
FNR	0.0329	0.0433
TNR	0.8890	0.0367
TPR	0.9671	0.0433
FPR	0.1109	0.0367

Table 2: Performance while training the network with emulated layer.




# Non-Differentiable Embedding Layer

## Invert Embedding Layer

Use a network to get the input sequence back from the embeddings.

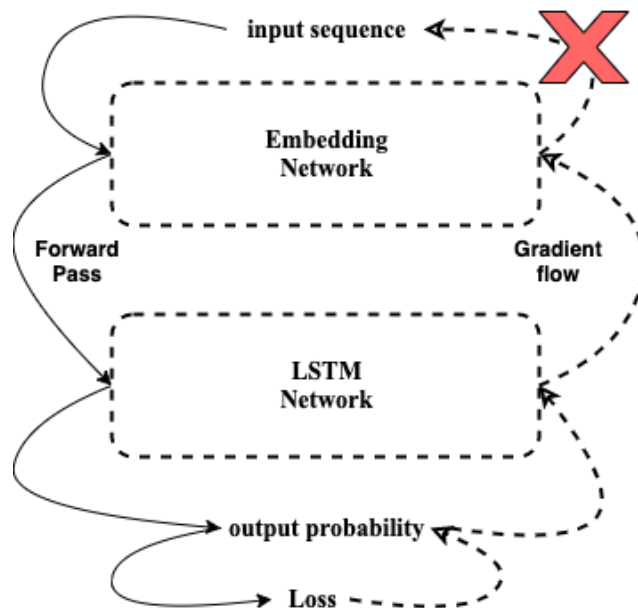
- How does that help us?
- This won't let gradients flow back to the input sequence.
- Let us see how...

# Non-Differentiable Embedding Layer



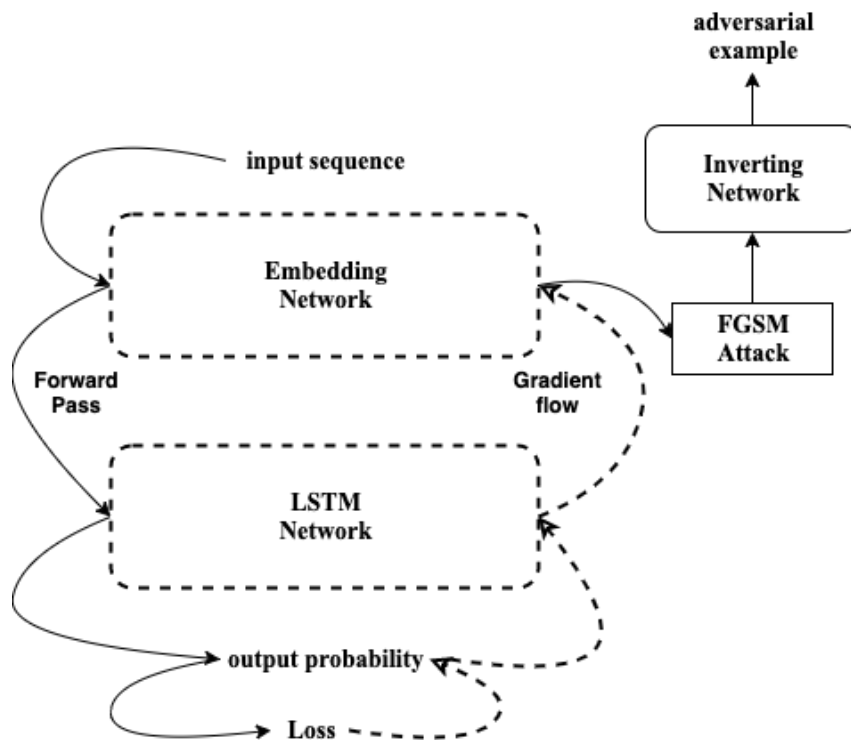
A forward pass through the network.

# Non-Differentiable Embedding Layer



**Our original setting.**

# Non-Differentiable Embedding Layer



**The inverting solution.**

# Non-Differentiable Embedding Layer

## Training the inverting network

*inputs* : Output of the Keras embedding layer.

*label* : The input sequence.

Trained similar to an autoencoder.

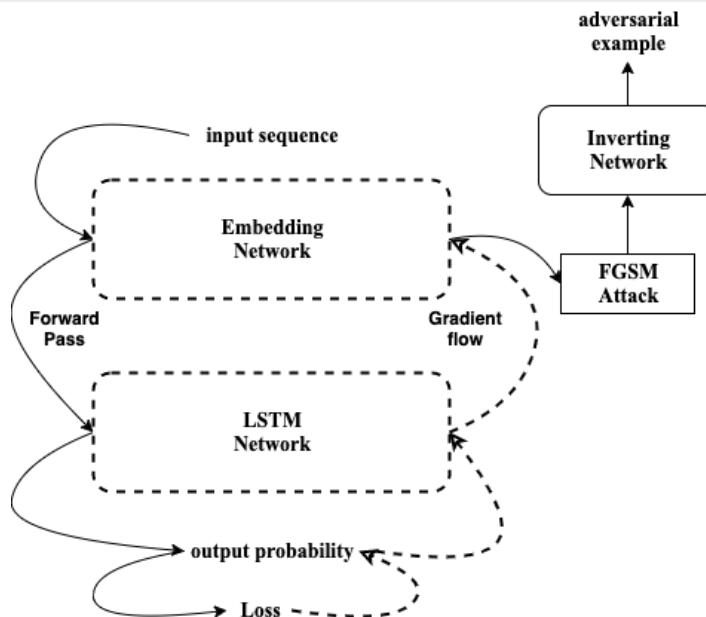
Input domain	Inverted domain
amazon.co.de	amazon.co.de
this-is-it-security.rwth	this-is-it-security.rwth
google.com	google.com
gst.gov.in	gst.gov.in
apple.com	apple.com

Table 3: Evaluation of the inverting network.

# Gradient-based Attack

## FGSM attack

$$x \leftarrow x + \epsilon \text{sgn}(\nabla_x L(\theta, x, y))$$



# Gradient-based Attack

## FGSM attack

$$x \leftarrow x + \epsilon \text{sgn}(\nabla_x L(\theta, x, y))$$

Keep attacking till probability drops below 10% or max-epochs reached.

## Generated adversarial samples

*-9.z07?\*hvd].a  
!n !fγ=s\*1d/.;  
jf/;;g;;0?2;*

- *Problem* : None of the adversarial domains are valid.
- *Catch* : The benign dataset was unfiltered.

# FGSM attack

- *Problem* : None of the adversarial domains are valid.
- *Catch* : The benign dataset was unfiltered.
- *Solution* : Remove invalid domains from the dataset. Retrain.

## Valid generated adversarial samples

*rjqqyq4q.net*  
*9cbq48qq.space*  
*wt34h8o0.space*  
*qqqqqqqq.net*


- *Catch* : Keep the top-level domain unchanged.
- *Too much change?* :  $x \leftarrow x + \epsilon \nabla_x L(\theta, x, y)$
- *Time/epoch* : 0.7092s



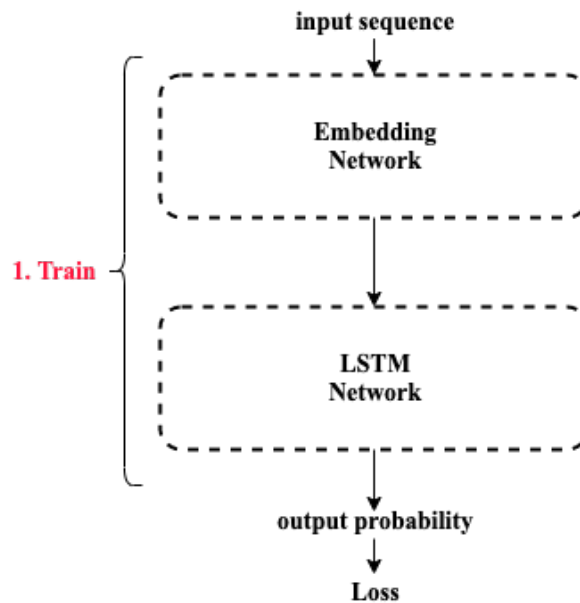
## The algorithm

- Train the network.
- *Inner Maximization* : Generate adversarial samples.
- Augment the training set with the adversarial samples.
- *Outer Minimization* : Retrain.


# Iterative Hardening



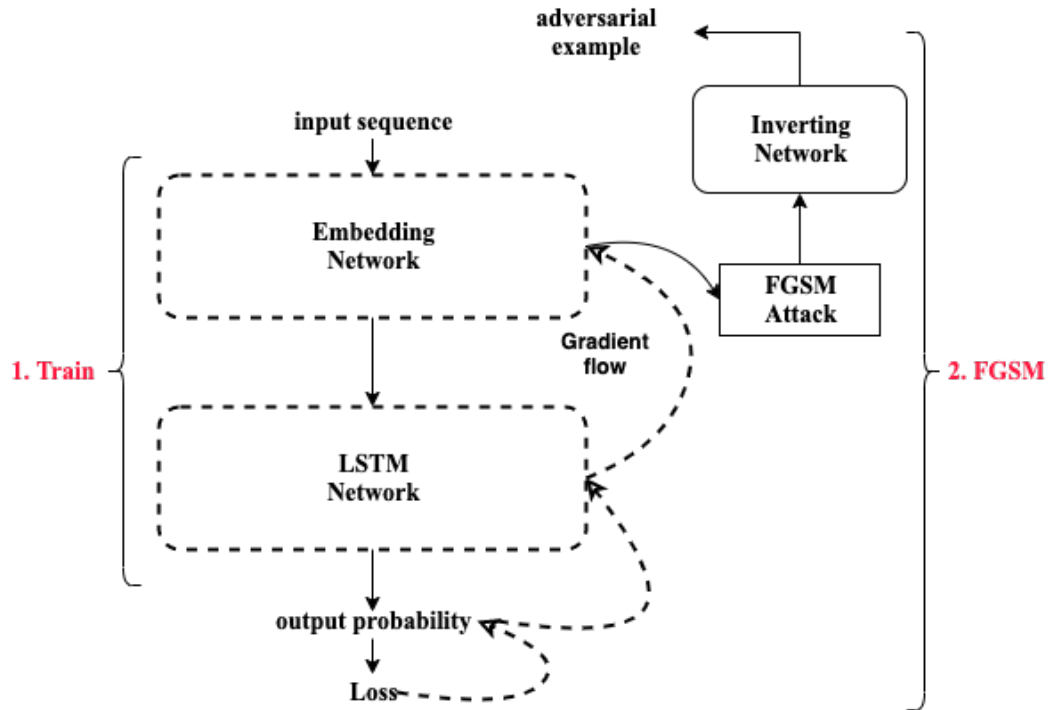
# Iterative Hardening



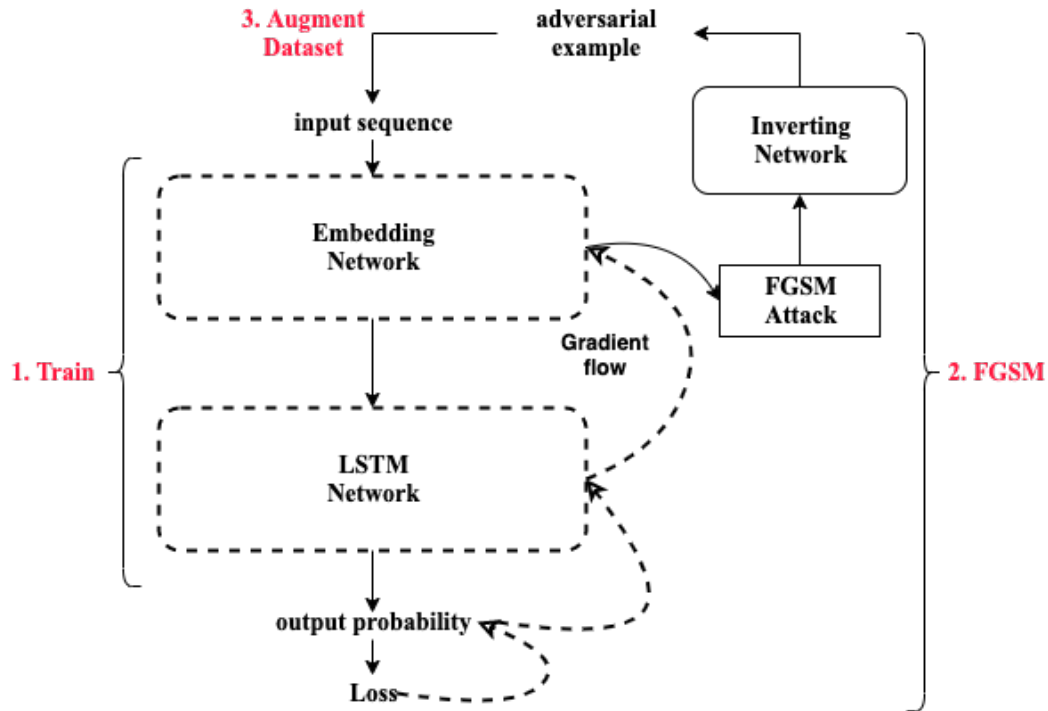
# Iterative Hardening



# Iterative Hardening



# Iterative Hardening



## The algorithm

- Train the network.
- *Inner Maximization* : Generate adversarial samples.
- Augment the training set with the adversarial samples.
- *Outer Minimization* : Retrain.
- *Expectation* : If it all works well, FGSM should generate **actually** benign samples.

Remember to freeze the embedding network after the first iteration.

## Results

057754af.space

[99.7%]



# Iterative Hardening

## Results

057754af.space

[99.7%]

FGSM


1.5	0.01	0.9	8.8
2.3	2.3	11.2	9.1
2.1	3.4	1.0	1.3
9.6	1.1	0.1	2.3

▪ ▪ ▪ ▪  
▪ ▪ ▪ ▪

[<5%]

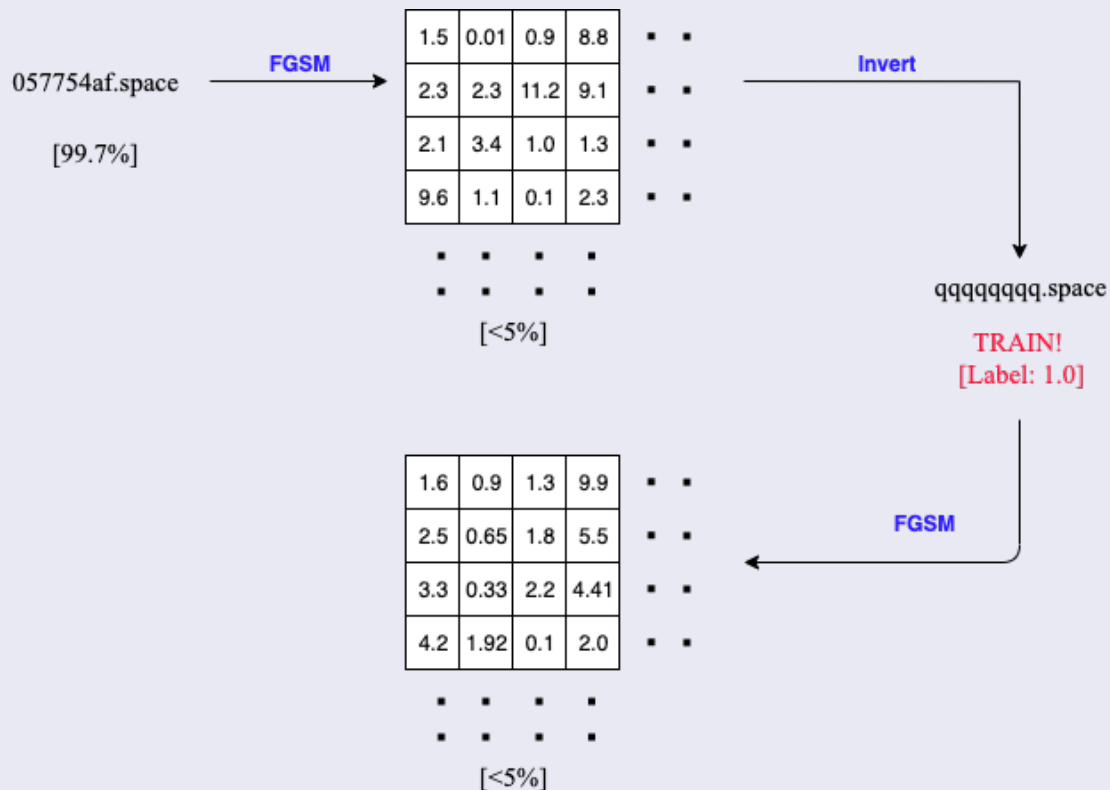
# Iterative Hardening

## Results



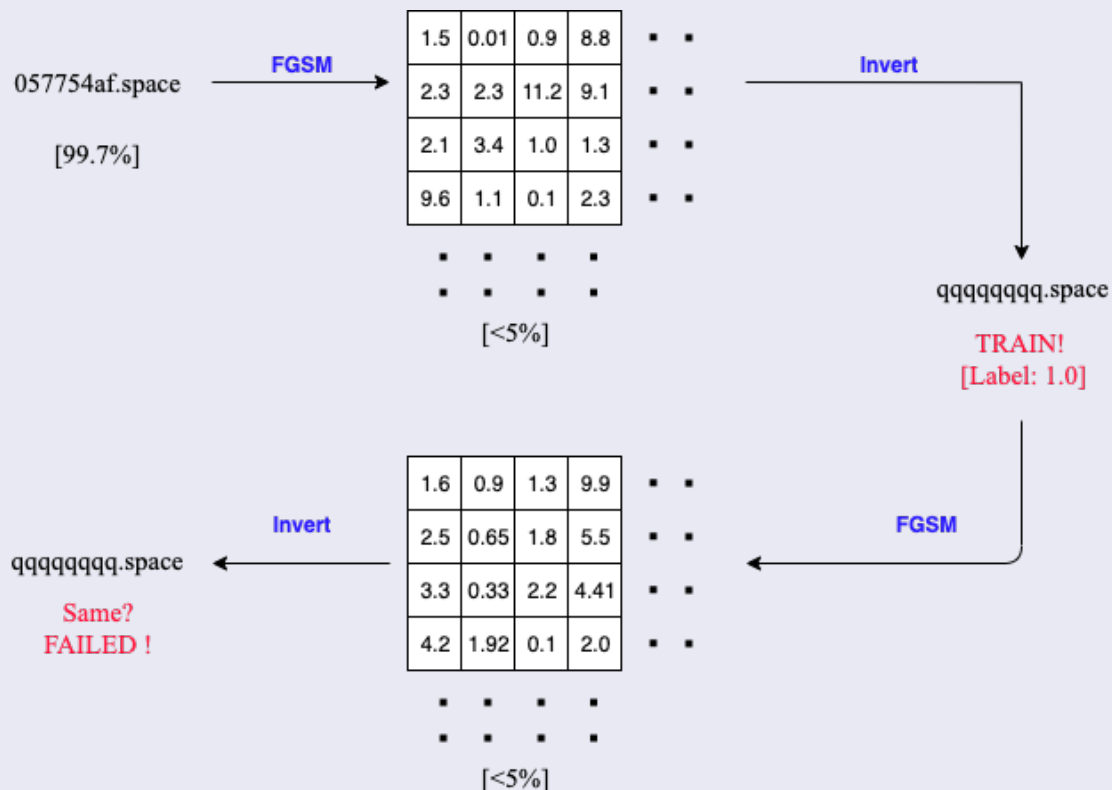
# Iterative Hardening

## Results



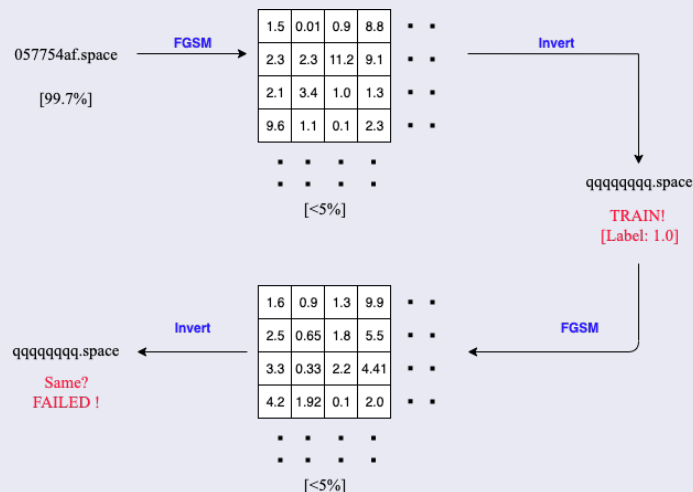
# Iterative Hardening

## Results



# Iterative Hardening

## Results



## Analysis

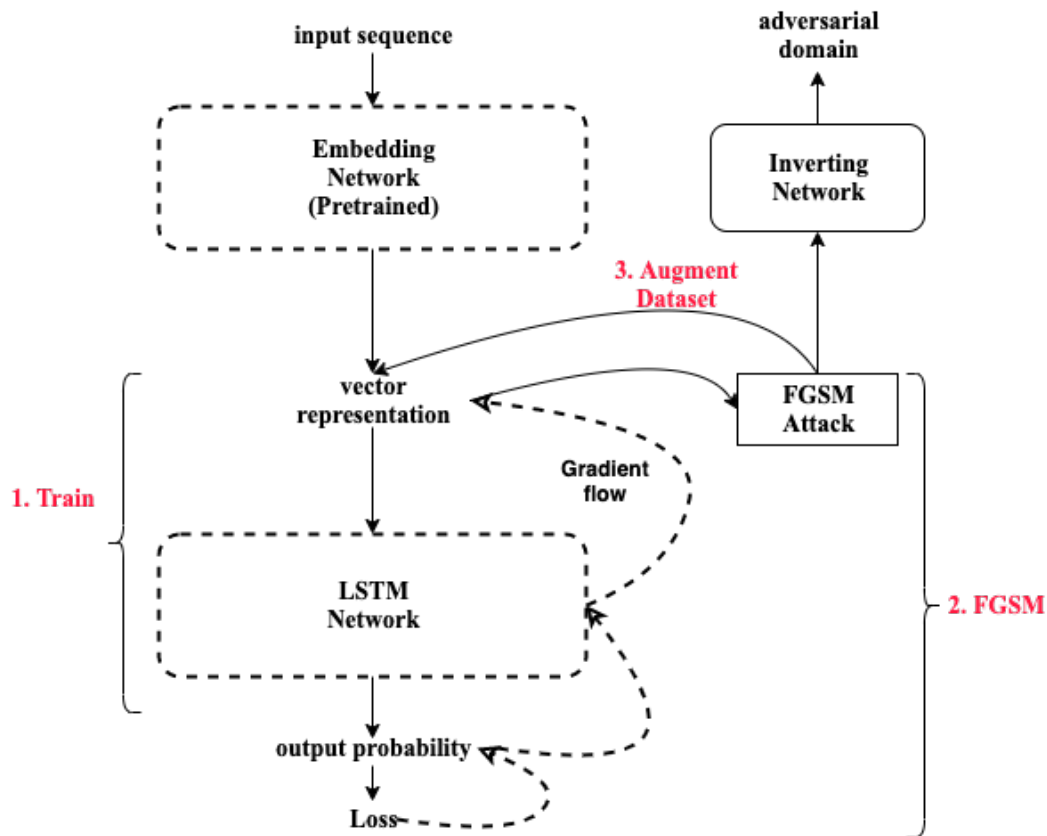
- *Vector Representation* :  $-0.07421777$   $0.05503434$   $0.10850348$  ...
- *Mean L2 Distance* :  $5.89923604974$
- *Baseline Prediction* : `qqqqqqqqq.space` `[99.8%]`

# Iterative Hardening over vector representations

## The algorithm

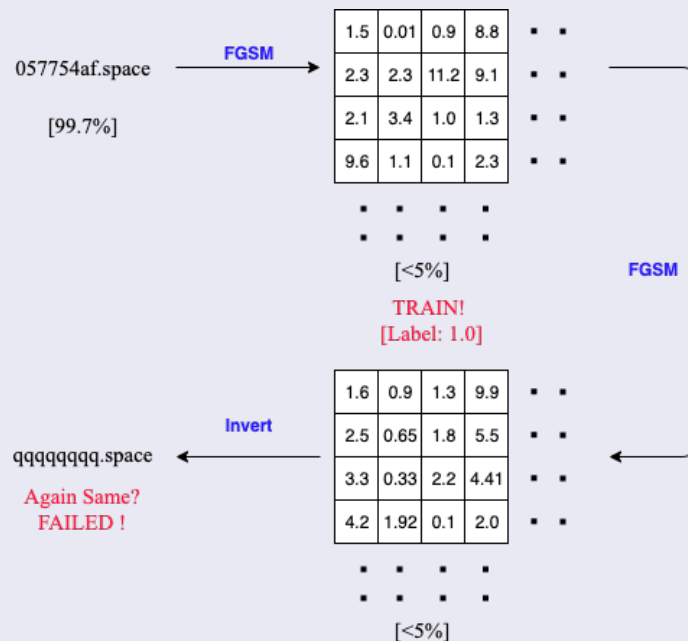
- Train the complete network.
- Create a dataset with vector representations.
- *Inner Maximization* : Generate adversarial vector representations.
- Augment the dataset with the adversarial samples.
- *Outer Minimization* : Retrain only the LSTM Network with this dataset.
- *Expectation* : If it all works well, FGSM should generate **actually** benign samples.

# Iterative Hardening over vector representations



# Iterative Hardening over vector representations

## Results



- *Average Epochs per Attack* : Increases on each iteration
- Harder to find adversaries.



## Why gradient based attacks don't work?

- The Embedding Layer **selects** a vector corresponding each character in the input sequence.
- *Discrete* : A character can be represented by a unique vector in the high dimensional space.
- The gradient based attack makes a continuous change in the direction of gradient.
- *Iterative Training in the character space* : Mapping the continuous change to discrete levels can disrupt the attack.
- *Iterative Training in the vector space* : The adversarial vector will never be generated by the embedding network. No useful learning.

# Transferring Adversaries

Mike Lorang, in his Master Thesis used transferred adversaries from a network without embedding layer to one with embedding layer.

-	Baseline	CharIterH	VectorIterH
Acc	0.9427	0.9739	0.9010
FNR	0.0573	0.0260	0.0989

Table 4: Evaluation against Transferred FGSM (LSTM).

-	Baseline	CharIterH	VectorIterH
Acc	0.9993	0.9879	0.9538
FNR	0.0007	0.0121	0.0462

Table 5: Evaluation against Transferred FGSM (CNN).

# Transferring Adversaries

-	Baseline	CharIterH	VectorIterH
Acc	0.7050	0.8860	0.7650
FNR	0.2950	0.1140	0.2350

Table 6: Evaluation against Hotflip Adversaries (LSTM).

-	Baseline	CharIterH	VectorIterH
Acc	0.5696	0.7264	0.4629
FNR	0.4303	0.2735	0.5370

Table 7: Evaluation against Hotflip Adversaries (CNN).

-	Baseline	CharIterH	VectorIterH
Acc	0.4667	0.6000	0.4667
FNR	0.5333	0.4000	0.5333

Table 8: Evaluation against SeqGAN Adversaries (Very few samples).

# Transferring Adversaries

Benign	Hotflip (CNN)
195.126.129.124.in-addr.arpa zjekmjf.germanistik.rwth-aachen.de ejgvgxp.ad.fh-aachen.de fe-prg007.nos-avg.cz	wli-hcg-.de z0n-e8tzmz7mrby-.be 5jkat2oz5gz8ei2.name zt-sf-lm.at

SeqGAN	Hotflip (LSTM)
qbutbtwbswul7a6anl.laanwh.ad e0ehl136oesqe.sfpspeeld.a.th qbutbtwbswul7a6anl.laanwh.ad o21pfr2o.e.s.hn	9-qqeidkufm28qd9j1.fr 9tnl777ld53b758.org e0dbmmgsm2-uav1-.jp bwh3pku3qm9e7.nz

Table 9: Benign and adversarial domain names.

# Future Work

## Flow gradients to the inputs without accuracy loss

- Train the original network.
- Generate embeddings
- Train the emulated embedding layer with the generated embeddings as labels.
- Improved the Emulated Embedding Layer?

## Projection of adversarial vector representation

- Generate adversarial vector representations.
- Calculate distance of the adversarial vectors to all possible embedding vectors. (L2 Distance)
- Choose the corresponding nearest embedding vectors.
- Projection from continuous to discrete space.

## Adversarial examples for discrete data / text

- SeqGAN: Sequence Generative Adversarial Nets with Policy Gradient (2017), Lantao Yu et al.
- HotFlip: White-Box Adversarial Examples for Text Classification (2018), Javid Ebrahimi et al.

Try other attacks for discrete data.